

# Comparative Analysis of Google Cloud Natural Language API and IBM Watson Natural Language Understanding for Advanced Text Analytics

Jinang Gandhi  
jinangg2@illinois.edu  
University of Illinois Urbana-Champaign  
Illinois, USA

Nikunj Agarwal  
nikunja2@illinois.edu  
University of Illinois Urbana-Champaign  
Illinois, USA

## ACM Reference Format:

Jinang Gandhi and Nikunj Agarwal. 2025. Comparative Analysis of Google Cloud Natural Language API and IBM Watson Natural Language Understanding for Advanced Text Analytics. In . ACM, New York, NY, USA, 3 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 SUMMARY

This technical review presents a comparative analysis of two cloud-based Natural Language Processing (NLP) services: Google Cloud Natural Language API [2] and IBM Watson Natural Language Understanding (NLU) [3], across three core tasks: sentiment analysis, entity recognition, and text classification. Using a diverse dataset comprising news articles, Reddit comments, and product reviews, we implemented a modular, reproducible framework to process and evaluate inputs through both APIs. While Google demonstrated finer granularity in sentiment scoring and broader entity recognition with detailed metadata, IBM provided more interpretable sentiment labels and attached sentiment at the entity level, proving useful for reputation tracking. In text classification, Google's taxonomy was broader and more specific, whereas IBM offered hierarchical categories in simpler terms. The results highlight key strengths and limitations of each platform, offering insights into their respective suitability for different NLP-driven applications.

## 2 INTRODUCTION

Natural Language Processing (NLP) has become a core component of many applications, ranging from customer feedback analysis to social media monitoring and news categorization. As cloud-based NLP services continue to evolve, developers and organizations are presented with multiple options to embed language understanding into their workflows. Two widely adopted platforms in this space are Google Cloud Natural Language API and IBM Watson Natural Language Understanding (NLU). While both offer similar capabilities such as sentiment analysis, entity recognition, and content classification, they differ in how they process and interpret language data, as well as in performance and ease of use.

---

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).  
*Conference'17, July 2017, Washington, DC, USA*

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.  
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM  
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

This project aims to conduct a side-by-side comparison of these two services by applying their APIs to a diverse set of real-world text inputs, including news articles, Reddit comments, and product reviews. Using custom-built scripts [1], we collected and structured this input data, then processed it through both the Google and IBM NLP pipelines. The outputs were stored and analyzed to understand how each service performs across key NLP tasks and input types. This report documents our findings, outlines our methodology, and discusses notable patterns, strengths, and limitations observed in each service.

## 3 DESCRIPTION

This project involved the design and implementation of a comparative framework to evaluate two widely used cloud-based Natural Language Processing (NLP) services: Google Cloud Natural Language API and IBM Watson Natural Language Understanding (NLU). The focus was on three core NLP tasks provided by both services: sentiment analysis, entity recognition, and content classification. The following subsections describe the data collection process, system architecture, implementation details, and output structure used for this evaluation.

### 3.1 Data Collection

To ensure diversity and realism in testing, the input data was sourced from three different categories of textual content:

- **Articles:** A set of recent BBC news articles were collected using the newspaper library, which allowed automated extraction of article text from provided URLs.
- **Comments:** Reddit comments were fetched using the PRAW (Python Reddit API Wrapper) library. We filtered results using keyword searches (e.g., “exceptional” in *r/science*).
- **Reviews:** To test sentiment interpretation specifically, product-related comments were fetched from subreddits like *r/phone* using relevant keywords such as “iPhone”.

The final dataset, stored in `input.json`, was structured with three top-level keys: `articles`, `comments`, and `reviews`, each holding a list of text entries for that category.

### 3.2 API Integration and Execution

Separate scripts were developed to interact with each NLP service. Each script loaded the input dataset, processed the text items by category, and stored the results in a structured JSON file for further analysis.

**3.2.1 Google Cloud Natural Language API.** The `gcp_api.py` script contains functions that send HTTP requests to three of Google's endpoints:

- **Sentiment Analysis:** Measures emotional tone using sentence-level and document-level polarity and magnitude scores.
- **Entity Recognition:** Identifies named entities (people, places, prices, organizations, etc.) along with metadata such as confidence and salience.
- **Text Classification:** Predicts content categories using Google's predefined taxonomy.

All responses are collected and stored under a common structure in `gcp_output.json`, preserving the original text, output type, and API responses.

**3.2.2 IBM Watson NLU.** A Jupyter Notebook (`ibm_nlu_api.ipynb`) was used to interact with IBM's NLU service. It calls equivalent endpoints for:

- **Sentiment Analysis:** Returns overall document sentiment and associated confidence.
- **Entity Recognition:** Includes entity type, relevance, count, confidence, and entity-level sentiment.
- **Category Classification:** Labels content using IBM's content categories and includes a confidence score.

The processed output is stored in `ibm_output.json` following a format similar to that used for Google results.

Both scripts were written with error handling in place to manage API failures, empty fields, or unsupported content. The results are stored incrementally in a list and finally serialized into a formatted JSON file using Python's `json` module.

### 3.3 Output Structure

Each output JSON file contains an array of result objects. Each object includes:

- A type field indicating the source type (article, comment, or review)
- The original text entry
- A nested sentiment dictionary
- A nested entities dictionary
- A nested categories dictionary

This structure allowed for easy comparison across APIs and categories.

For example, in the Google output, sentiment values include both magnitude and score for each sentence, while IBM returns a labeled polarity (positive, neutral, or negative) with a confidence score. In entity recognition, Google groups entities by type and provides metadata such as Wikipedia links or numeric values (for PRICE, DATE, etc.). IBM adds fine-grained sentiment for each entity, including contextual polarity.

### 3.4 Automation and Reproducibility

The overall pipeline was designed to run as a batch process. Once the input data is available, running the scripts triggers the full analysis cycle: from sending API requests to generating side-by-side comparable output files. The code is modular and can be reused with different datasets or expanded to include more NLP tasks in the future.

## 4 EVALUATION

To evaluate the performance and capabilities of Google Cloud Natural Language API and IBM Watson NLU, we processed identical input data consisting of articles, Reddit comments, and product reviews, through both services. This section compares their outputs across three primary NLP tasks: sentiment analysis, entity recognition, and text classification. We highlight both the strengths and limitations observed in practice.

### 4.1 Sentiment Analysis

Both Google Cloud Natural Language API and IBM Watson NLU support sentiment analysis at the document level, but they differ significantly in output structure, interpretability, and granularity. Below are the detailed points of comparison:

- **Output Format:**
  - **Google:** Returns a score (range: -1 to 1) indicating sentiment polarity and a magnitude representing emotional intensity, irrespective of polarity.
  - **IBM:** Provides a labeled sentiment (positive, negative, or neutral), a continuous score, and an optional mixed flag indicating the presence of conflicting sentiments.
- **Granularity:**
  - **Google:** Supports sentence-level sentiment scoring. This granularity is beneficial for analyzing longer texts with varied tones.
  - **IBM:** Typically returns a single sentiment label for the entire document. Sentence-level sentiment is not as detailed or consistently applied.
- **Performance in Practice:**
  - In a news article combining neutral reporting with negative political commentary:
    - \* **Google:** Successfully highlighted negative sentiment in political references while maintaining neutrality in factual parts.
    - \* **IBM:** Returned a single sentiment label that overlooked these nuances.
  - In a Reddit comment expressing both praise and criticism:
    - \* **Google:** Assigned mixed sentiment scores across different sentences, accurately capturing tonal shifts.
    - \* **IBM:** Labeled the entire comment as neutral, missing the emotional variation.
- **Interpretability:**
  - **IBM:** More accessible for non-technical users due to its labeled outputs, which are easy to understand at a glance.
  - **Google:** While more detailed, its numeric scores require more interpretation or visualization to extract meaningful insights.
- **Mixed Sentiment Handling:**
  - **IBM:** Includes a mixed flag, but its usage was inconsistent and didn't always reflect shifts in tone within the document.
  - **Google:** Does not have an explicit mixed flag but inherently captures mixed sentiment through sentence-level granularity and score variation.

Overall, Google offers a more granular and analytically rich sentiment output, while IBM prioritizes interpretability and ease of

use. Depending on the use case, technical analysis vs. high-level summary, either service can be advantageous.

## 4.2 Entity Recognition

Both Google Cloud NLP and IBM Watson NLU offer strong entity recognition capabilities with complementary strengths. Google provides broader and more granular entity coverage, while IBM excels in sentiment-aware entity analysis. Below are some key observations from our evaluation:

- **Google Cloud NLP:**
  - Identified a wide range of entity types such as *dates*, *prices*, *locations*, *organizations*, and *persons*.
  - Returned *metadata* like Wikipedia links and normalized values for entities.
  - Supported *multiple mentions* of the same entity, capturing both *contextual relevance* and *frequency*.
  - Offered *fine-grained categorization* (e.g., distinguishing between Organization, Location, Person).
  - Included *mention offsets*, aiding downstream tasks like co-reference resolution or information extraction.
  - **Example:** In a Reddit comment about Apple, Google correctly identified “Apple” as an Organization.
- **IBM Watson NLU:**
  - Focused on extracting *high-relevance named entities*.
  - Provided *sentiment scores for each entity*, which Google does not natively offer.
  - **Example:** In an article about DHL, IBM identified “DHL” and “\$800” with *negative sentiment*, reflecting the article’s context of trade restrictions.
  - Occasionally *generalized entity types* (e.g., labeling many as Organization).
  - Sometimes *missed less prominent entities*.
  - **Example:** In the same Reddit comment, IBM *did not return any entity*, unlike Google.

In summary, Google’s API offers richer entity data and broader recognition, while IBM adds unique value by incorporating sentiment analysis at the entity level—making both tools effective depending on the application context.

## 4.3 Category Classification

- **IBM Watson NLU**
  - Provides hierarchical category classification with associated confidence scores.
  - Example: For a DHL article, correctly identified topics like /law, govt and politics and /finance/financial news.
  - Taxonomy is somewhat limited in breadth and may omit low-confidence categories entirely.
  - Leans towards context-based classification, favoring a broader narrative understanding.
  - Returns a relatively flat JSON output, making it easier to parse and integrate in downstream applications.
- **Google Cloud NLP**
  - Uses its proprietary taxonomy, which is broader and often returns more specific categories.

- Performs better on longer, well-structured text inputs; for example, assigns detailed categories to comprehensive news articles.
- For shorter or less coherent text inputs (e.g., Reddit comments), often fails to return categories.
- May produce unexpected or overly general classifications – e.g., a tech article on AI regulation was categorized under Computers & Electronics, missing policy-related nuances.
- Favors topical interpretation over narrative or contextual nuance.
- Returns a nested JSON structure that may require additional parsing logic for integration and analysis.

## 5 DISCUSSION

**Overall API Comments:** Google Cloud NLP excels in structured, detailed analysis and is well-suited for sentence-level parsing, entity linking, and tasks requiring precision and control. However, its output is dense and may be less intuitive for non-technical audiences.

IBM Watson NLU provides cleaner, more accessible outputs, with sentiment tied to entities and categories that often align well with human interpretation. However, it may overlook finer details or return fewer entities in complex inputs.

**Ease of Integration:** Both Google Cloud NLP and IBM Watson NLU offer a seamless integration experience. With just an account setup and API key generation, developers can quickly start using the REST APIs. While Google uses API keys and IBM uses service credentials, the overall process is straightforward and well-documented in both cases, making them equally accessible for developers.

**Documentation Quality:** Both platforms provide comprehensive and easy-to-navigate documentation. Google’s documentation is detailed with summaries, request/response formats, and interactive API testing directly in the browser. Similarly, IBM Watson NLU offers clear descriptions, example requests, and thorough parameter explanations. Although their presentation styles differ, both deliver high-quality resources that effectively support developers throughout the integration and development process.

Ultimately, both APIs are capable, and their utility depends on the application. For use cases involving document annotation, advanced analytics, or fine-grained sentiment tracking, Google is a stronger choice. For summarization dashboards or use by non-experts, IBM offers a more user-friendly interface and interpretation.

## REFERENCES

- [1] Jinang Gandhi and Nikunj Agarwal. 2025. NLP API Analysis: A Comparative Study of Google Cloud NLP and IBM Watson NLU. [https://github.com/nikunjagarwal321/nlp\\_api\\_analysis](https://github.com/nikunjagarwal321/nlp_api_analysis). Accessed: 2025-05-02.
- [2] Google Cloud. 2024. Google Cloud Natural Language API. <https://cloud.google.com/natural-language>. <https://cloud.google.com/natural-language>.
- [3] IBM Cloud. 2024. IBM Watson Natural Language Understanding. <https://www.ibm.com/cloud/watson-natural-language-understanding>. <https://www.ibm.com/cloud/watson-natural-language-understanding>.